

Semaphores and Related System Calls

Apurba Sarkar

IIST Shibpur

February 18, 2019

- 1 Semaphore
 - Motivation
 - Few important structures
 - System calls needed

What are they actually?

- Semaphores are system-implemented data structures that are shared between processes. Some features of the semaphore are
 - 1 Semaphores are used to synchronize operations when processes access a common, limited, and possibly non-shareable resource.
 - 2 Each time a process wants to obtain the resource, the associated semaphore is tested. A positive, non-zero semaphore value indicates the resource is available.
 - 3 Semaphore system calls will, by default, cause the invoking process to block if the semaphore value indicates the resource is not available.

Uses of Semaphores

- Guarding Critical Section
 - Semaphores that control access to a single resource, taking the value of 0 (resource is in use) or 1 (resource is available), are often called binary semaphores.
- Precedence Enforcer
 - Semaphores ensuring an ordering of execution among concurrent processes
- Resource Counter
 - Semaphores controlling access to N (multiple) resources, thus assuming a range of non-negative values, are frequently called counting semaphores.

Things to be discussed

- Creating and Accessing Semaphore Sets
 - Semaphores are allocated in sets.
 - They are allocated with a value of 0 (blocked)
- Semaphore Operations
 - wait (P)
 - signal (V)
- Semaphore Control
 - get info on a set
 - set all semaphore values in a set or get all values
 - set one value in a set or get one value
 - remove the set

Structures to be used

- semid_ds

```
struct  semid_ds {  
    . . .  
    struct  sem  *sem_base; /*pointer to first semaphore in set*/  
    ushort_t  sem_nsems; /* number of semaphores in set */  
};
```

structures to be used

- Structure sem

```
struct    sem {  
    ushort    semval; /*semaphore value*/  
    pid_t      sempid; /*pid of last operation */  
    ushort    semcnt; /* # waiting for increase in semval */  
    ushort    semzcnt; /* # waiting for semval == 0 */  
};
```

System Call `semget`

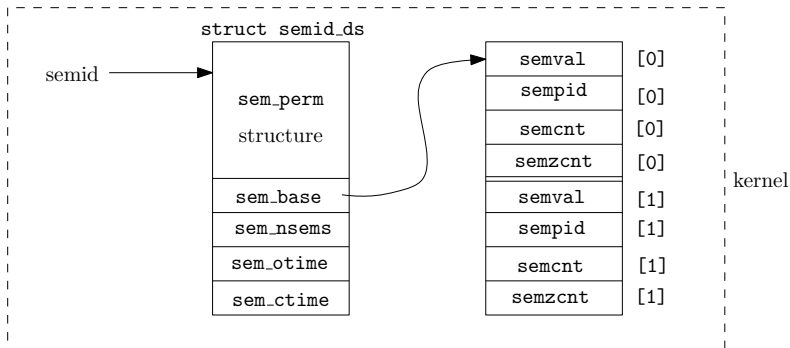
- Function of `semget`
 - To create a semaphore, or gain access to one that exists.
- to include `<sys/types.h>` `<sys/ipc.h>` `<sys/sem.h>`
- System Call: `int semget(key_t key, int nsems, int semflg);`
- Return values:
 - **On Success:** the semaphore identifier (`semid`)
 - **On failure:** `-1` and Sets `errno`
- Arguments
 - `key_t key`: used to identify a semaphore set
 - `int nsems`: the number of semaphores in the set.
 - `int semflg`: specify access permissions and/or special creation condition(s).

semget notes

- If the value for key does not have a semaphore identifier associated with it, and `IPC_CREAT` has been specified, a new set of semaphores is created (`semget`).
- If key is `IPC_PRIVATE`, and `IPC_CREAT` has been specified, a new set of semaphores is created (`semget`).
- If `IPC_CREAT` is specified (but not `IPC_EXCL`), and the semaphore set for the indicated key value already exists, the (`semget`) call will return the associated semaphore identifier.
- When using `semget` to access an established semaphore set, the value of `nsems` can be set to 0 (a dont-care value).
- If a semaphore identifier exists for this key and `IPC_CREAT` and `IPC_EXCL` are both set, (`semget`) will fail and the “file exists” error message will be returned via `errno`.

semget

- We can picture a particular semaphore in the kernel as being a `semid_ds` structure that points to an array of `sem` structure.



semop Call - Semaphore Operation

- Function of `semop`
 - to perform operations on individual semaphores..
- to include `<sys/types.h>` `<sys/ipc.h>` `<sys/sem.h>`
- System Call: `int semop (int semid, struct sembuf *sops, size_t nsops);`
- Return values:
 - On Success: 0
 - On failure: -1 and Sets `errno`
- Arguments
 - `int semid`: semaphore identifier.
 - `struct sembuf *sops`: a reference to the address of an array of semaphore operations that will be performed on the semaphore set denoted by the `semid` value.
 - `size_t nsops`: the number of elements in the array of semaphore operations.

Semaphore Buffer for `semop` Call

```
struct    sembuf {  
    ushort    sem_num;           /* semaphore # */  
    short     sem_op;           /* semaphore operation */  
    short     sem_flg;          /* operation flags */  
};
```

- `ushort sem_num`: semaphore number (the index into the array of sem structures referenced by the `sem_base`).
- `short sem_op`: operation to be performed on the semaphore.
- `short sem_flg`:
 - `IPC_NOWAIT`: if the semaphore operation can not be performed, the call will return immediately.
 - `SEM_UNDO`: allows an operation to be undone if a blocked operation subsequently fails.
 - `0`: process blocks until semaphore available

Actions Taken by `semop`

If `sem_op` value:

- **is positive:** Add `sem_op` to `semval`. This is a release of a resource.
- **is zero:** The caller will block until the semaphores value becomes zero.
- **is negative:** The caller is blocked until the semaphores value (`semval`) becomes greater than or equal to the absolute value of `sem_op`. Then, the absolute value of `sem_op` is subtracted from `semval`.

semctl System Call - Semaphore Control

- Function

- To perform a variety of generalized control operations on the system semaphore structure, on the semaphores as a set and on individual semaphores.

- to include `<sys/types.h>` `<sys/ipc.h>` `<sys/sem.h>`

- System Call: `int semctl(int semid, int semnum, int cmd, /* union semun arg*/ ...);`

- Return values:

- On Success: 0
- On failure: -1 and Sets `errno`

- Arguments

- `int semid`: a valid semaphore identifier.
- `semnum`: the number of semaphores in the semaphore set.
- `int cmd`: an integer command value (`IPC_STAT`, `IPC_SET`, ..).
- `arg`: union of type `semun`.

cmd values of `semctl` Call

- `IPC_STAT`: return the current values of the `semid_ds` structure for the indicated semaphore identifier.
- `IPC_SET`: modify a restricted number of members in the `semid_ds` structure.
- `IPC_RMID`: remove the semaphore set.
- `GETALL`: return the current values of the semaphore set.
- `SETALL`: initialize all semaphores in a set to the values stored in the array referenced by the fourth arguments to `semctl`.
- `GETVAL`: return the current of the individual semaphore referenced by the value of the `semnum` argument.
- `SETVAL`: set the value of a single semaphore in a set

union semun in the semctl call

- A union is is a later-day version of the Pascal variant `record`. It is a data structure that can take on multiple forms.
- `semctl()` requires a union to handle the different kinds of data that can be provided to it or received from it.

```
union  semun {  
    int          val;  
    struct semid_ds *buf;  
    ushort       * array;  
} arg; // declares a semun named arg
```

- The value in `arg` is one of:
 - `int val`: an integer (0 or others),
 - `struct semid_ds *buf`: a reference to a `semid_ds` structure,
 - `ushort *array`: the base address of an array of short integers (the values for the semaphore(s)).